


Qwen3-0.6B, all 196 decoder linears quantized

(a) Distortion vs. rate



(b) Perplexity vs. rate

